



NTNU – Trondheim
Norwegian University of
Science and Technology

Visualization of large scale Netflow data

Nicolai Eeg-Larsen

Submission date: May 2016
Responsible professor: Yuming Jiang, ITEM
Supervisor: Otto Wittner, UNINETT

Norwegian University of Science and Technology
Department of Telematics

Title: Visualization of large scale Netflow data
Student: Nicolai Eeg-Larsen

Problem description:

Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

This is the second paragraph. Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

Responsible professor: Yuming Jiang, ITEM
Supervisor: Otto Wittner, UNINETT

Sammendrag

Sikkerheten til nesten all offentlig nøkkel-kryptografi er basert på et vanskelig beregnbarhetsproblem. Mest velkjent er problemene med å faktorisere heltall i sine primtallsfaktorer, og å beregne diskrete logaritmer i endelige sykliske grupper. I de to siste tiårene, har det imidlertid dukket opp en rekke andre offentlig nøkkel-systemer, som baserer sin sikkerhet på helt andre type problemer. Et lovende forslag, er å basere sikkerheten på vanskeligheten av å løse store likningssett av flervariabe polynomlikninger. En stor utfordring ved å designe slike offentlig nøkkel-systemer, er å integrere en effektiv “falluke” (trapdoor) inn i likningssettet. En ny tilnærming til dette problemet ble nylig foreslått av Gligoroski m.f., hvor de benytter konseptet om kvasigruppe-strengtransformasjoner (quasigroup string transformations). I denne masteroppgaven beskriver vi en metodikk for å identifisere sterke og svake nøkler i det nylig foreslalte multivariable offentlig nøkkel-signatursystemet MQQ-SIG, som er basert på denne idéen.

Vi har gjennomført et stort antall eksperimenter, basert på Gröbner basis angrep, for å klassifisere de ulike parametrene som bestemmer nøklen i MQQ-SIG. Våre funn viser at det er store forskjeller i viktigheten av disse parametrene. Metodikken består i en klassifisering av de forskjellige parametrene i systemet, i tillegg til en innføring av konkrete kriterier for hvilke nøkler som bør velges. Videre, har vi identifisert et unødvendig krav i den originale spesifikasjonen, som krevde at kvasigruppene måtte oppfylle et bestemt kriterie. Ved å fjerne denne betingelsen, kan nøkkelerings-algoritmen potensielt øke ytelsen med en stor faktor. Basert på alt dette, foreslår vi en ny og forbedret nøkkel-genereringsalgoritme for MQQ-SIG, som vil generere sterkere nøkler og være mer effektiv enn den originale nøkkel-genereringsalgoritmen.

Preface

Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

Acknowledgements

Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

Contents

List of Figures	ix
Listings	xi
List of Acronyms	xiii
1 Background	1
1.1 NetFlow	1
1.1.1 How does it work?	1
1.1.2 Main components	3
1.1.3 nfdump	3
1.2 Data visualization	5
1.2.1 Characteristics	6
1.2.2 Visual perception	6
1.2.3 Data presentation architecture	7
1.3 D3.js	8
1.3.1 How does it work?	8
2 Research	13
2.1 Related work	13
2.2 Initial research	14
2.3 Traits of a DDoS attack	15
2.3.1 Raw NetFlow format	16
3 D3.js and NetFlow	19
3.1 Using D3.js	19
3.1.1 Scope	19
3.2 Number of flows to a certain host and port	21
3.2.1 Scope in D3.js	21
4 Challenges	25
4.1 Large data sets	25
4.1.1 IP-spectrum	25

4.1.2 Increasing number of flows	25
References	27
Appendices	
A Appendix A	31
B Appendix B	41

List of Figures

1.1	Creating a flow in the NetFlow cache siter	2
1.2	Figure of a simple NetFlow architecture	3
1.3	Example of dataset of random numbers where no pre-attentive processing is done	4
1.4	Example of dataset of random numbers where no pre-attentive processing is done	4
1.5	Example of dataset of random numbers where no pre-attentive processing is done	5
1.6	Example of dataset of random numbers where no pre-attentive processing is done	7
1.7	Example of a dataset of random numbers where pre-attentive processing has been used to distinguish the occurrences of the number five	7
2.1	Ten files in the provided files with the most flows [cis16]	15
2.2	Ten files in the provided files with the most flows [cis16]	16
2.3	Ten files in the provided files with the most flows	17
3.1	Ten files in the provided files with the most flows	20
3.2	The smallest files from the provided files	20
3.3	Top ten used destination addresses within the timeframe 1300-1400, 18th of January	20
3.4	Top ten used destination addresses within the timeframe 1300-1400, 18th of January	21
3.5	Top ten used destination addresses within the timeframe 1300-1400, 18th of January	22
3.6	Top ten used destination addresses within the timeframe 1300-1400, 18th of January	23
3.7	Top ten used destination addresses within the timeframe 1300-1400, 18th of January	23

Listings

1.1	HTML example	8
-----	------------------------	---

List of Acronyms

AS Autonomous Systems.

BGP Border Gateway Protocol.

CSS Cascading Style Sheets.

CSV Comma Separated Value.

DDoS Distributed Denial of Service.

DoS Denial of Service.

DPA Data Presentation Architecture.

HTML HyperText Markup Language.

IP Internet Protocol.

IPFIX IP Flow Information eXport.

IPv4 Internet Protocol version 4.

IPv6 Internet Protocol version 6.

LVM Logical Volume Manager.

NTNU Norwegian University of Science and Technology.

SVG Scalable Vector Graphics.

TPC Transmission Control Protocol.

Chapter 1

Background

1.1 NetFlow

Cisco IOS NetFlow creates an environment that have the tools to understand who, what, when, where and how network traffic is flowing. This makes it easier for administrators to utilize the network as optimal as possible. One can determine the source and destination of traffic and use this information to reveal for example DDoS-attacks or spam mail.

1.1.1 How does it work?

Every packet that is forwarded within a router/switch is examined for a set of Internet Protocol (IP) packet attributes. With these attributes one can determine if the packet is unique or similar to other packets.

The attributes used by NetFlow are:

- IP source address
- IP destination address
- Source port
- Destination port
- Layer 3 protocol type
- Class of service
- Router/Switch interface

To group packets into a flow, one compares source/destination IP address, source/destination ports, protocol interface and class of service. Then the packets

2 1. BACKGROUND

and bytes are tallied. This method is scalable because a large amount of network information is condensed into a database of NetFlow information called the NetFlow cache.

When the NetFlow cache is created one can use this to understand the network behaviour. The different attributes generate different knowledge about a certain network, and combined they can paint a detailed picture of how the network is working. For example the ports show what application is utilizing the traffic, while the tallied packets and bytes show the amount of traffic. [SST⁺16]

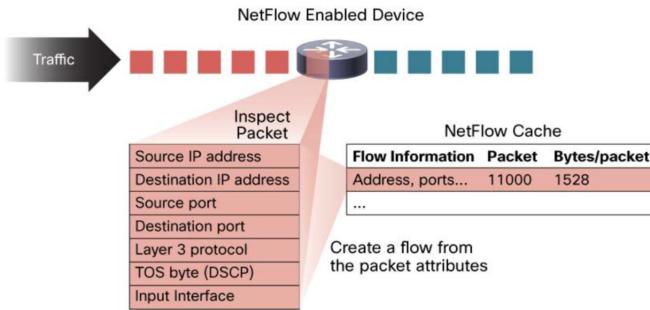


Figure 1.1: Creating a flow in the NetFlow cache siter

- Source address allows the understanding of who is originating the traffic
- Destination address tells who is receiving the traffic
- Ports characterize the application utilizing the traffic
- Class of service examines the priority of the traffic
- The device interface tells how traffic is being utilized by the network device
- Tallied packets and bytes show the amount of traffic

Additional information added to a flow includes:

- Flow timestamps to understand the life of a flow; timestamps are useful for calculating packets and bytes per second
- Next hop IP addresses including Border Gateway Protocol (BGP) routing Autonomous Systems (AS)
- Subnet mask for the source and destination addresses to calculate prefixes

- flags to examine Transmission Control Protocol (TCP) handshakes

1.1.2 Main components

A typical set-up using NetFlow consists of three main components:

- **Flow Exporter:** aggregates packets into flows and exports flow records towards one or more flow collectors.
- **Flow collector:** is responsible for reception, storage and pre-processing of flow data received from a flow exporter.
- **Analysis application:** an application that analyze the received flow data in different contexts, such as intrusion or traffic profiling.

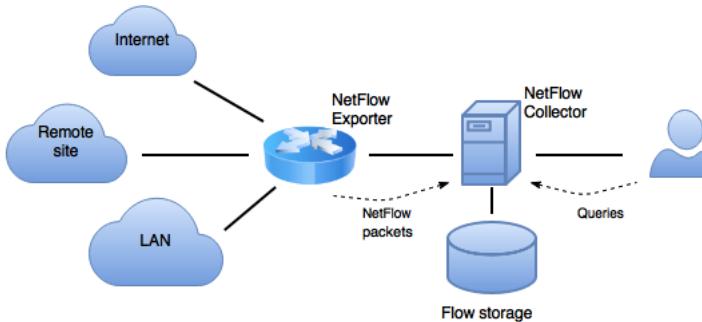


Figure 1.2: Figure of a simple NetFlow architecture

1.1.3 nfdump

skrive om ipfix, v5 og v9, muligens i egen seksjon) nfdump collect and process NetFlow data on the command line. It stores NetFlow data in time sliced files. The files are binary and this provides the possibility of either returning the output from nfdump in the same binary form, or as readable text. nfdump has four output formats, raw, line, long and extended. The challenge of representing Internet Protocol version 6 (IPv6) addresses is handled by shrinking them in the normal output. In figure 1.3 the collection process is depicted, and in figure 1.4 the processing of collected NetFlow data is shown.[nfd16]

Lime inn
hvordan
det ser
ut i kom-
mandolin-
jen

sitere
listen

4 1. BACKGROUND

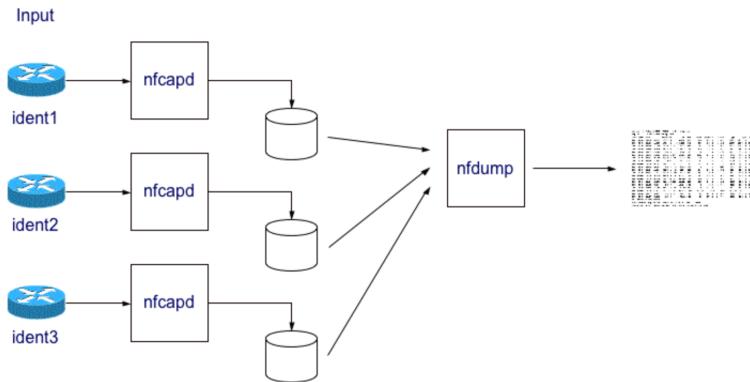


Figure 1.3: Example of dataset of random numbers where no pre-attentive processing is done

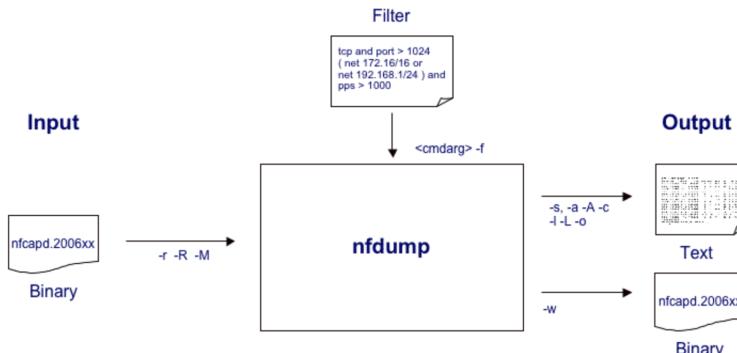


Figure 1.4: Example of dataset of random numbers where no pre-attentive processing is done

v5,v9 and IPfix

- **v5:** NetFlow v5 is definitely the most popular version of Cisco Netflow. It is fixed, meaning it always stays the same and makes for a simpler deciphering.
- **v9:** v9 is opposite of its predecessor dynamic. The collector will need to know the format of incoming NetFlow v9 flows, which means v9 templates periodically needs to be sent to the collector to inform of the format which the flows are being exported are. It was made to support technologies as Multi-cast, IPSec and Multi Protocol Label Switching (MPLS). This thanks to the templates. IPv6 support was added as well.

- **IPFIX:** Based on the design of NetFlow v9, IP Flow Information eXport (IPFIX) added support for variable length strings. Making it possible for Application Visibility and Control(AVC) exports in the future.

forklare?

Example of use

An example of a nfdump command used in this project is for example how to extract the number of flows each day, and find the 10 most used destination IP-addresses:

```
1 nfdump -R /data/netflow/oslo_gw/2012/01/01/nfcapd.201201010000:nfcapd
  .201201012355 -n 10 -s dstip -o csv > example.csv
```

Such a request iterates over a number of files due to the -R command. In this case it is all captures between 00:00 until 23:55 on the first of January 2012. It is limited to the 10 most popular destination IP's with the -n and -s. All of this is stored in a .Comma Separated Value (CSV)file which is optimal for use with the D3.js framework.

This results in the output shown in figure 1.5.

Date first seen	Duration Proto	Dst IP Addr	Flows(%)	Packets(%)	Bytes(%)	pps	bps	bpp
2011-12-31 23:58:52.073	86420.475 any	191.229.233.80	304171(2.3)	374411(0.4)	369.4 M(0.9)	4	34194	986
2011-12-31 23:58:50.706	86421.658 any	162.185.32.85	204693(1.5)	210870(0.2)	29.1 M(0.1)	2	2693	137
2011-12-31 23:58:50.872	86416.387 any	161.220.8.250	171799(1.3)	175078(0.2)	7.8 M(0.0)	2	721	44
2011-12-31 23:58:50.748	86421.701 any	162.185.32.105	171177(1.3)	367501(0.4)	21.7 M(0.1)	4	2009	59
2011-12-31 23:58:50.178	86422.648 any	196.49.180.97	137249(1.0)	380318(0.4)	296.5 M(0.7)	4	27447	779
2011-12-31 23:58:52.698	86419.411 any	161.223.1.186	116698(0.9)	158786(0.2)	25.4 M(0.1)	1	2355	160
2011-12-31 23:58:50.762	86421.892 any	159.152.49.239	114587(0.9)	696163(0.8)	953.8 M(2.3)	8	88288	1370
2011-12-31 23:58:52.421	86419.823 any	190.49.138.164	98869(0.7)	110728(0.1)	6.8 M(0.0)	1	632	61
2011-12-31 23:58:52.881	86419.483 any	190.49.138.168	84728(0.6)	109403(0.1)	161.8 M(0.4)	1	14977	1478
2012-01-01 15:19:13.044	29561.745 any	71.238.74.19	84632(0.6)	592943(0.7)	793.9 M(1.9)	20	214838	1338
Summary: total flows: 13373369, total bytes: 41.8 G, total packets: 86.5 M, avg bps: 3.9 M, avg pps: 1000, avg bpp: 482								
Time window: <unknown>								
Total flows processed: 13373369, Blocks skipped: 0, Bytes read: 695426912								
Sys: 5.373s flows/second: 2488661.8 Wall: 5.407s flows/second: 2473155.2								

Figure 1.5: Example of dataset of random numbers where no pre-attentive processing is done

1.2 Data visualization

Data visualization refers to the techniques used to communicate data or information by encoding it as visual objects[sitere?]. Meaning that information is represented through any visual element such as graphs and plots, but may also take any other visual form. Visualization helps users analyse and interact with data in a whole new way. It makes complex data more accessible, understandable and usable.[Fri08]

In recent years the rate of which data is generated has increased rapidly, and the need for information to be available and comprehensible is growing. All these new sources of data has created what we refer to as "Big Data". Without visual

6 1. BACKGROUND

presentation such data is too big to understand. This is the big reason for visualization is emerging as a big market.

Combining several parameters through visualization could reveal something automated systems might ignore or don't pick up on.

The greatest value of a picture is when it forces us to notice what we never expected to see.

by John Tukey.

1.2.1 Characteristics

In his book from 1983, *The Visual Display of Quantitative Information*[TGM83], Edward Tufte defines characteristics any effective graphical representation should contain as:

- show the data
- induce the viewer to think about the substance rather than about methodology, graphic design, the technology of graphic production or something else
- avoid distorting what the data has to say
- present many numbers in a small space
- make large data sets coherent
- encourage the eye to compare different pieces of data
- reveal the data at several levels of detail, from a broad overview to the fine structure
- serve a reasonably clear purpose: description, exploration, tabulation or decoration
- be closely integrated with the statistical and verbal descriptions of a data set.

1.2.2 Visual perception

In this paper the correlation between effective visual communication and how it is perceived upon human inspection is important. A humans ability to distinguish between differences in length, shape and color is referred to as "pre-attentive attributes".

A good example of this is imagining finding the number of a certain character in a series of characters. This requires significant time and effort, but if the character were to stand out by being a different size, color or orientation this could be done quickly through pre-attentive processing. Good data visualization takes all of this into consideration and uses pre-attentive processing. In this simple example it is easy to see how pre attentive processing is used to distinguish how many occurrences of the number 5 is in a larger set of random numbers.

```
987349790275647902894728624092406037070570279072
803208029007302501270237008374082078720272007083
247802602703793775709707377970667462097094702780
927979709723097230979592750927279798734972608027
```

Figure 1.6: Example of dataset of random numbers where no pre-attentive processing is done

```
987349790275647902894728624092406037070570279072
803208029007302501270237008374082078720272007083
247802602703793775709707377970667462097094702780
927979709723097230979592750927279798734972608027
```

Figure 1.7: Example of a dataset of random numbers where pre-attentive processing has been used to distinguish the occurrences of the number five

1.2.3 Data presentation architecture

Data Presentation Architecture (DPA) has its purpose to identify, locate, manipulate, format and present data in such a way as to optimally communicate meaning and proffer knowledge[wik16]. This has become an important tool in Business Intelligence, the art of transforming raw data into something useful.

Objectives

DPA has two main objectives, which is the following:

8 1. BACKGROUND

- To use data to provide knowledge in the most efficient manner possible (minimize noise, complexity, and unnecessary data or detail given each audience's needs and roles)
- To use data to provide knowledge in the most effective manner possible (provide relevant, timely and complete data to each audience member in a clear and understandable manner that conveys important meaning, is actionable and can affect understanding, behaviour and decisions)

Scope

The actual work of DPA consist of:

- Creating effective delivery mechanisms
- Define relevant knowledge needed by each viewer
- Determine how often the data should be updated
- Determine how often and when the user needs to see the data
- Finding the right data
- Utilizing the best visualizations and presentation formats

1.3 D3.js

In this paper D3.js [Bos12] is chosen as the framework to create examples of effective data visualizations due to its dynamical and interactive properties. D3 stands for Data-Driven Documents, and is a Javascript library. D3.js allows users to bind arbitrary data to a Document Object Model. It uses widely implemented Scalable Vector Graphics (SVG), Cascading Style Sheets (CSS) and HyperText Markup Language (HTML)5 standards. D3 is unique in the way it creates SVG objects from large datasets using simple D3.js functions to generate rich text/graphic charts and diagrams.

1.3.1 How does it work?

The W3C DOM API is often tiring to use. An example bit of code from[link/kilde] shows how one changes the text color of paragraph elements:

```
1 var paragraphs = document.getElementsByTagName( "p" );
3 for (var i = 0; i < paragraphs.length; i++) {
    var paragraph = paragraphs.item(i);
    paragraph.style.setProperty( "color" , "white" , null );
```

```
| }
```

Listing 1.1: HTML example

In D3.js this could be solved through one line of code:

```
8 d3.selectAll("p").style("color", "white");
```

Listing 1.2: D3.js example

D3.js also possess dynamic properties which gives the user a powerful tool to create advanced graphics with a small amount of code.

This next snippet of code shows how the D3.js framework simply appends to an existing html object.

```

<!DOCTYPE html>
10 <meta charset="utf-8">
11 <style> /* set the CSS */
12   body { font: 12px Arial; }
13
14   path {
15     stroke: steelblue;
16     stroke-width: 2;
17     fill: none;
18   }
19
20   .axis path,
21   .axis line {
22     fill: none;
23     stroke: grey;
24     stroke-width: 1;
25     shape-rendering: crispEdges;
26   }
27
28 </style>
29 <body>
30
31
32 <!-- load the d3.js library -->
33 <script src="http://d3js.org/d3.v3.min.js"></script>
34
35 <script>
36
37   // Set the dimensions of the canvas / graph
38   var margin = {top: 30, right: 20, bottom: 30, left: 50},
39     width = 600 - margin.left - margin.right,
40     height = 270 - margin.top - margin.bottom;
41
42   // Parse the date / time
43   var parseDate = d3.time.format("%d-%b-%y").parse;
```

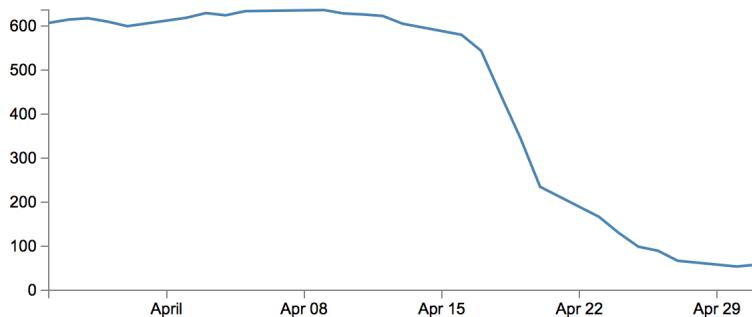
10 1. BACKGROUND

```
44 // Set the ranges
45 var x = d3.time.scale().range([0, width]);
46 var y = d3.scale.linear().range([height, 0]);
47
48 // Define the axes
49 var xAxis = d3.svg.axis().scale(x)
50   .orient("bottom").ticks(5);
51
52 var yAxis = d3.svg.axis().scale(y)
53   .orient("left").ticks(5);
54
55 // Define the line
56 var valueline = d3.svg.line()
57   .x(function(d) { return x(d.date); })
58   .y(function(d) { return y(d.close); });
59
60 // Adds the svg canvas
61 var svg = d3.select("body")
62   .append("svg")
63     .attr("width", width + margin.left + margin.right)
64     .attr("height", height + margin.top + margin.bottom)
65   .append("g")
66     .attr("transform",
67       "translate(" + margin.left + ", " + margin.top + ")");
68
69 // Get the data
70 d3.csv("data.csv", function(error, data) {
71   data.forEach(function(d) {
72     d.date = parseDate(d.date);
73     d.close = +d.close;
74   });
75
76   // Scale the range of the data
77   x.domain(d3.extent(data, function(d) { return d.date; }));
78   y.domain([0, d3.max(data, function(d) { return d.close; })]);
79
80   // Add the valueline path.
81   svg.append("path")
82     .attr("class", "line")
83     .attr("d", valueline(data));
84
85   // Add the X Axis
86   svg.append("g")
87     .attr("class", "x_axis")
88     .attr("transform", "translate(0, " + height + ")")
89     .call(xAxis);
90
91   // Add the Y Axis
92   svg.append("g")
93     .attr("class", "y_axis")
94     .call(yAxis);
```

```
96 }));  
98 </script>  
100 </body>
```

Listing 1.3: Example of use of the D3.js framework

legg til kilde på koden her. This graph would be appended to the body element of the html and look like this:



In the code the dynamic properties are visible as the x- and y-axis change its parameters based on the input data.

Chapter 2

Research

2.1 Related work

In the last decade the importance of security against attacks on large computer systems has grown rapidly. In 2004, the ACM workshop on Visualization and data mining for computer security presented NVisionIP: netflow visualizations of system state for security situational awareness[LYBL04]. This was one of the first tools to visualize NetFlow data. The visualization was based on either number of bytes transmitted or the number of flows to or from the hosts on the network.

In [LYL04] they discuss the use of NVisionIP to combat different security concerns. Most of the same attacks covered in this paper are relevant today, only in today's massive amounts of data, they may be way more difficult to discover.

- **Worm infection:** One of the most basic security function one might uncover. Worms usually spread by probing for other hosts. Filtering out hosts transmitting a lot of Flows with a single destination port, one could easily see which machines are infected and should be taken offline.
- **Compromised systems:** If a host is compromised, the attacker might install malware that allows the attacker to control the machine. Following this an attacker might turn a host into a file server. By detecting large volumes of traffic on certain ports one might discover such an attack.
- **Misuse:** Misuse of computer networks in order with terms of use etc.. An example is detecting if certain users have abnormal high volumes of traffic, and by inspecting in more detail one can uncover if this through one single application and not in accordance with the policies of the organization.
- **Port Scans:** When a large number of ports are used at a specific host it is easily identified by NVisionIP.

- **Denial of Service (DoS):** Denial of Service Attacks will be visible through spikes in traffic volume from the host attacking. If a host is attacked the same pattern is visible through high volumes in receiving traffic. Thus peaks in traffic is not necessarily an attack, but might be a result of a new release, or backup etc.

2.2 Initial research

In section 1.1 we see how the raw format of the NetFlow packets look. Norwegian University of Science and Technology (NTNU) Logical Volume Manager (LVM) Comparing how understandable this format is comparing to a visual representation will be the main object of this paper(omformulere. ikke helt riktig). How much more effective is visualization compared to the raw format read by machines.

To understand this, experiments will determine how quickly one can distinguish an attack from both the raw format, and the visual representation.

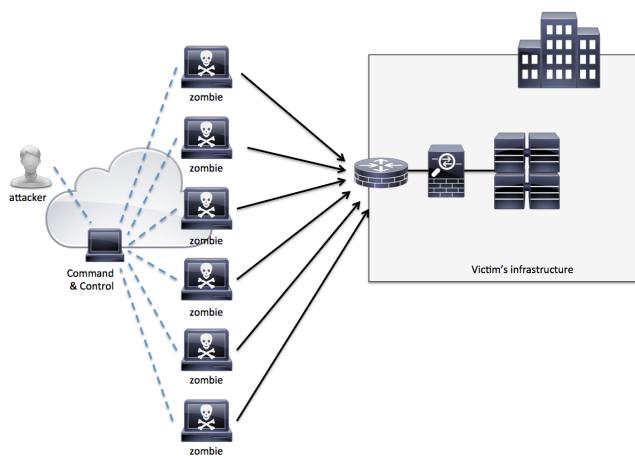
This is where D3.js will come to great use. It can be used to quickly develop simple interactive graphs that can be used to test theories up against each other.

To be able to identify an DDoS attack, one can look at it from two angles. By finding someone whom is attacking, or someone whom is being attacked. In this case we will look at the second scenario. As mentioned earlier simply a peak in flows is not enough grounds to establish an actual attack. First of all, one will need to look for patterns of similar incidents, and what lies behind them.

2.3 Traits of a DDoS attack

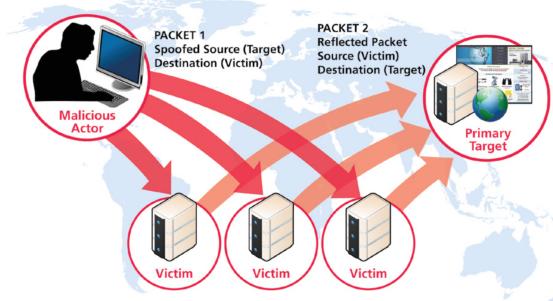
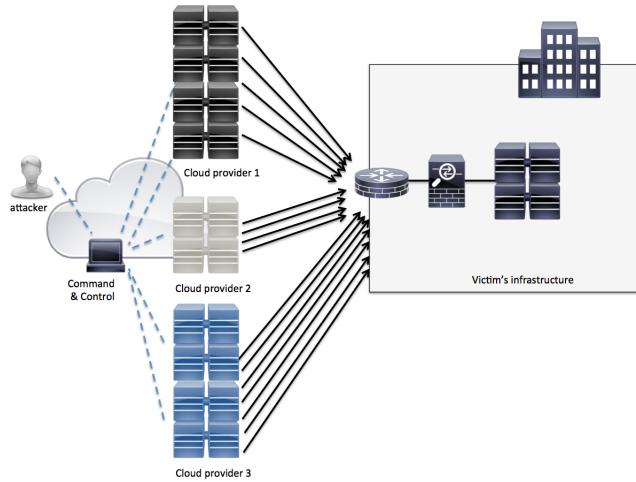
In a Distributed Denial of Service (DDoS) attack there are a large number of hosts performing the attack. In many cases a lot of them are not even aware they are a part on an attack. This is called a botnet, derived from the words robot and network. Using compromised systems, called zombies, gives the attacker control of a large enough amount of hosts to perform a volume-based DDoS attack.

Figure 2.1: Ten files in the provided files with the most flows [cis16]



Another new trend that has emerged is using large datacenters or cloud machines to launch these attacks. Either through renting or compromising them. As cloud providers are offering such large amounts of computers, this new platform is not only great for legitimate use, but also cyber-criminals.

Distributed Reflection Denial of Service attacks are becoming more and more popular. DrDoS techniques usually involve multiple victim host machines that unwillingly participate in a DDoS attack on the attacker's primary target. Requests to the victim host machines are redirected, or reflected, from the victim hosts to the target. Anonymity is one advantage of the DrDoS attack method. In a DrDoS attack, the primary target appears to be directly attacked by the victim host servers, not the actual attacker. This approach is called spoofing. Amplification is another advantage of the DrDoS attack method. By involving multiple victim servers, the attacker's initial request yields a response that is larger than what was sent, thus increasing the attack bandwidth.

Figure 2.2: Ten files in the provided files with the most flows [cis16]

2.3.1 Raw NetFlow format

In the preceding example, there are multiple flows for UDP port 80 (hex value 0050). In addition, there are also flows for TCP port 53 (hex value 0035) and TCP port 80 (hex value 0050).

The packets in these flows may be spoofed and may indicate an attempt to perform these attacks. It is advisable to compare the flows for TCP port 53 (hex value 0035) and TCP port 80 (hex value 0050) to normal baselines to aid in determining whether an attack is in progress.

```

router#show ip cache flow
IP packet size distribution (90784136 total packets):
 1-32 64 96 128 160 192 224 256 288 320 352 384 416 448 480
 .000 .698 .011 .001 .004 .005 .000 .004 .000 .000 .003 .000 .000 .000 .000
 512 544 576 1024 1536 2048 2560 3072 3584 4096 4608
 .000 .001 .256 .000 .010 .000 .000 .000 .000 .000 .000 .000 .000 .000 .000
IP Flow Switching Cache, 4456704 bytes
 1885 active, 63651 inactive, 59960004 added
 129803821 ager polls, 0 flow alloc failures
 Active flows timeout in 30 minutes
 Inactive flows timeout in 15 seconds
IP Sub Flow Cache, 402056 bytes
 0 active, 16384 inactive, 0 added, 0 added to flow
 0 alloc failures, 0 force free
 1 chunk, 1 chunk added
 last clearing of statistics never
Protocol      Total    Flows   Packets  Bytes  Packets  Active(Sec)  Idle(Sec)
-----      Flows     /Sec    /Flow   /Pkt   /Sec    /Flow     /Flow
TCP-Telnet    11393421    2.8      1    48     3.1     0.0     1.4
TCP-FTP       236        0.0      12    66     0.0     1.8     4.8
TCP-FTPD      21         0.0    13726   1294    0.0    18.4     4.1
TCP-WWW       22282        0.0     21   1020    0.1     4.1     7.3
TCP-X          719        0.0      1    40     0.0     0.0     1.3
TCP-BGP        1         0.0      1    40     0.0     0.0    15.0
TCP-Frag      70399        0.0      1    688     0.0     0.0    22.7
TCP-other     47861004   11.8      1   211    18.9     0.0     1.3
UDP-DNS       582        0.0      4    73     0.0     3.4    15.4
UDP-NTP       287252        0.0      1    76     0.0     0.0    15.5
UDP-other     310347        0.0      2   230     0.1     0.6    15.9
ICMP          11674        0.0      3    61     0.0    19.8    15.5
IPv6INIP      15         0.0      1  1132     0.0     0.0    15.4
GRE           4         0.0      1    48     0.0     0.0    15.3
Total:      59957957   14.8      1   196    22.5     0.0     1.5
SrcIf      SrcIPAddress   DstIf      DstIPAddress   Pr SrcP DstP   Pkts
G10/0      192.168.10.201 Gi0/1      192.168.60.102 06 0984 0050   1
G10/0      192.168.11.54  Gi0/1      192.168.60.158 06 0911 0035   3
G10/1      192.168.150.60 Gi0/0      10.89.16.226 06 0016 12CA   1
G10/0      192.168.10.17  Gi0/1      192.168.60.97 11 0B89 0050   1
G10/0      10.88.226.1   Gi0/1      192.168.202.22 11 007B 007B   1
G10/0      192.168.12.185 Gi0/1      192.168.60.239 11 0BD7 0050   1
G10/0      10.89.16.226  Gi0/1      192.168.150.60 06 12CA 0016   1
router#

```

Figure 2.3: Ten files in the provided files with the most flows

Chapter 3

D3.js and NetFlow

3.1 Using D3.js

Earlier in this paper it is mentioned that D3.js will be used to show examples of effective visualization of NetFlow data. It is assumed that the data has already been processed before it is made accessible to these examples. I was supplied with two months of anonymous data from UNINETT to get familiar with NetFlow and be able to use real data for my visualizations. This is anonymized data from January of 2012 from Trondheim and Oslo NetFlow collectors. This means millions of flows.

3.1.1 Scope

The vast amounts of data should be presented with such a scope that is intuitive and easily understandable. Considering NetFlow packages is timestamped and sent from one source address to a specific destination address's port, one will have to chose which of these spectrum's to focus on. In the visual solution it is natural to combine these to represent the data.

IP spectrum

Choosing the address spectrum as the main focus, one will have to find a way to represent the entire IPv4 spectrum. This is alone a challenge, and when it comes to IPv6, it becomes practically impossible. This results in relying more on the pre-processing of the data and segregating the IP-addresses actually worth noticing. In the data provided by UNINETT it is possible to list for example the top 10 files in size, meaning more flows. In the data provided by UNINETT this search provided the results in figure 3.1.

From this simple preprocessing it is easy to see that in the time period between 1300-1400 on the 18th of January there was a clear peak in the number of flows having all the spots in the top 10. If we compare to the times with the lowest amount there is a different as they are a fraction of the others.

```
[eeglarse@iou2:/data/netflow$ find . -printf '%s %p\n'|sort -nr|head
14838848 ./oslo_gw/2012/01/18/nfcapd.201201181325
14729440 ./oslo_gw/2012/01/18/nfcapd.201201181335
14729284 ./oslo_gw/2012/01/18/nfcapd.201201181310
14720548 ./oslo_gw/2012/01/18/nfcapd.201201181330
14687944 ./oslo_gw/2012/01/18/nfcapd.201201181315
14651908 ./oslo_gw/2012/01/18/nfcapd.201201181340
14566420 ./oslo_gw/2012/01/18/nfcapd.201201181320
14563196 ./oslo_gw/2012/01/18/nfcapd.201201181305
14508804 ./oslo_gw/2012/01/18/nfcapd.201201181345
14472664 ./oslo_gw/2012/01/18/nfcapd.201201181300]
```

Figure 3.1: Ten files in the provided files with the most flows

```
[eeglarse@iou2:/data/netflow$ find . -printf '%s %p\n'|sort -nr|tail -100
849408 ./trd_gw1/2012/01/01/nfcapd.201201010920
848160 ./trd_gw1/2012/01/01/nfcapd.201201010745
842856 ./trd_gw1/2012/01/01/nfcapd.201201010725
834212 ./trd_gw1/2012/01/01/nfcapd.201201010655
832340 ./trd_gw1/2012/01/01/nfcapd.201201010640
830364 ./trd_gw1/2012/01/01/nfcapd.201201010555
828856 ./trd_gw1/2012/01/01/nfcapd.201201010845
821940 ./trd_gw1/2012/01/01/nfcapd.201201010900
816012 ./trd_gw1/2012/01/01/nfcapd.201201010905
804624 ./trd_gw1/2012/01/01/nfcapd.201201010735
802596 ./trd_gw1/2012/01/01/nfcapd.201201010545
799164 ./trd_gw1/2012/01/01/nfcapd.201201010635
780600 ./trd_gw1/2012/01/01/nfcapd.201201010650
772644 ./trd_gw1/2012/01/01/nfcapd.201201010610
760424 ./trd_gw1/2012/01/01/nfcapd.201201010705
758864 ./trd_gw1/2012/01/01/nfcapd.201201010720]
```

Figure 3.2: The smallest files from the provided files

```
[eeglarse@iou2:~$ cat test_180112.csv |cut -f 5 -d ',' |sort|uniq -c|sort|tail
17762 162.185.32.85
19878 161.222.192.123
21506 191.220.233.80
23995 161.223.1.164
37704 161.223.1.108
39759 159.152.145.176
49316 161.223.1.142
51467 190.49.180.97
61424 161.223.1.106
120976 192.239.62.2
```

Figure 3.3: Top ten used destination addresses within the timeframe 1300-1400, 18th of January

Through this we create a .csv file containing the hour in question going further in detail. Analysing which destination address is the most requested is the next step.

Again one specific address is clearly separated from the others. At this point we have gotten such into detail on the dataset, it is time to find the reason behind the results we have found. These high numbers could be a DDoS attack, or other

```
eeglarse@iou2:~$ cat test_180112.csv |cut -f 4 -d ',' |sort|uniq -c|sort|tail -10
18502 161.222.192.123
18557 191.220.233.80
19338 162.185.32.85
29367 161.223.1.164
46376 192.239.62.2
47139 190.49.180.97
47844 161.223.1.108
50509 161.223.1.142
77527 159.152.145.176
83184 161.223.1.106
```

Figure 3.4: Top ten used destination addresses within the timeframe 1300-1400, 18th of January

types of attacks, but not does not necessary ill willed. If we look at the list of top IP-addresses sending packets.

We see that the same IP-address, 192.239.62.2, is here high up as well. Among hundreds of thousands of addresses in the spectrum.

To further investigate the activity on certain IP-addresses, it is possible to get the most popular ports on either one specific IP-address, or a list.

sjekke
opp med
portnr og

Time spectrum

On the other side we have the time spectrum. In this case one looks at the amounts of flows within one time slot. Not down to the different IP-addresses. With the vast amount of IP-adresses this is not a suitable spectrum to present the data to find specific attacks etc.. On the other hand it could be used to monitor amounts of traffic over time or which ports are in use at certain times.

3.2 Number of flows to a certain host and port

This example shows how a simple graph can recognize a DDoS attack trough giving the option to see the number of netflows on different hosts and ports.

Rette
opp

3.2.1 Scope in D3.js

In this section three modules of a solution is presented to show different levels of detail. It combines both the time spectrum and IP-spectrum to investigate the NetFlow data in different ways. The data from UNINETT required pre-processing before being made available to the D3.js solution. The bash script used can be found in AppendixB.

File structure

With the bash scripts, thousands of files are created. Data is split into as small and many files as possible to reduce loading time at the user side, and to make sure the data is as comprehensible as possible.

Overview

First we have the overview which in this case show an entire year separated into months, weeks and days. The purpose of this is to be able to quickly recognize patterns in the data that correlates to periodically activities as backup or launches of new software as mentioned in 2.1. For example a weekly backup will create similar levels of data usage at specific times each week.

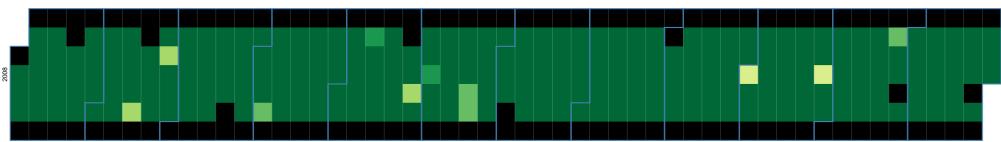


Figure 3.5: Top ten used destination addresses within the timeframe 1300-1400, 18th of January

IP-addresses and ports

For each day there are millions of different combinations of which IP-addresses and ports that send flows between each other. Through pre-processing it is possible to distinguish which IP-addresses are the most popular each day, and thus find the ports they use the most. This visualization shows the number of flows for each of these combinations through a heatmap. A heatmap means it distinguishes values through a color range based on the highest values in the data set.

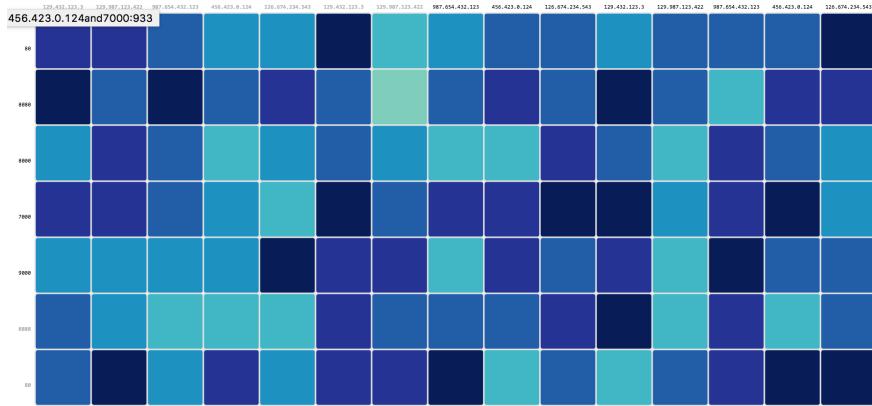


Figure 3.6: Top ten used destination addresses within the timeframe 1300-1400, 18th of January

24-hour chart

When a IP-address and port is selected for a specific day, the next scope is to look at the data in more detail. Using the time-spectrum this graph shows the 24-hour lapse and the amount of flows at each time.

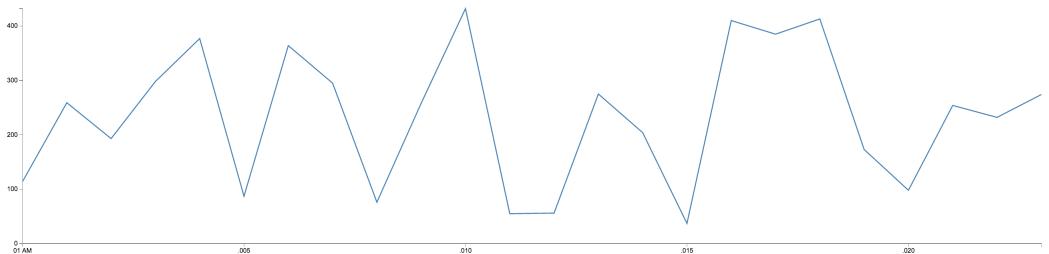


Figure 3.7: Top ten used destination addresses within the timeframe 1300-1400, 18th of January

Chapter 4 Challenges

4.1 Large data sets

When visualizing big data the main challenge is to effectively show the core message of the data. Considering one hour of the data provided from UNINETT, there is almost 400,000 different IP-adresses. And the amounts of flows is in the millions.

In section 1.2.1 good visualization is said to be able to present many numbers in a small space, make large data sets coherent, and reveal data at several levels of detail. I chose to create individual modules with D3.js, with each covering a different layer of detail.

4.1.1 IP-spectrum

As mentioned the range of the Internet Protocol version 4 (IPv4) is large, and with the emergence of IPv6 there is a challenge present. In 3.2.1 this was resolved with pre-processing of the data for a specific task. In other cases such a limitation on the number of IP-addresses represented wouldn't be satisfying.

4.1.2 Increasing number of flows

The amount of data sent these days are expanding quickly. This means the number of flows will follow, and a visual solution will need to be scalable to handle this increase. In the solution in 3.2

References

- [Bos12] Michael Bostock. D3. js. *Data Driven Documents*, 2012.
- [cis16] 2016.
- [Fri08] Vitaly Friedman. Data visualization and infographics smashing magazine, 2008.
- [LYBL04] K. Lakkaraju, W. Yurcik, R. Bearavolu, and A. J. Lee. Nvisionip: an interactive network flow visualization tool for security. In *Systems, Man and Cybernetics, 2004 IEEE International Conference on*, volume 3, pages 2675–2680 vol.3, Oct 2004.
- [LYL04] Kiran Lakkaraju, William Yurcik, and Adam J. Lee. Nvisionip: Netflow visualizations of system state for security situational awareness. In *Proceedings of the 2004 ACM Workshop on Visualization and Data Mining for Computer Security, VizSEC/DMSEC '04*, pages 65–72, New York, NY, USA, 2004. ACM.
- [nfd16] 2016.
- [SST⁺16] Products Services, Cisco Software, Cisco Technologies, Management Instrumentation, Cisco NetFlow, Data Literature, and White Papers. Introduction to cisco ios netflow - a technical overview, 2016.
- [TGM83] Edward R Tufte and PR Graves-Morris. *The visual display of quantitative information*, volume 2. Graphics press Cheshire, CT, 1983.
- [wik16] Data presentation architecture, 2016.

Todo list

Lime inn hvordan det ser ut i kommandolinjen	3
sitere listen	3
(.	3
forklare?	5
sjekke opp med portnr og	21
Rette opp	21

Appendix A

```
1  <!DOCTYPE html>
3  <meta charset="utf-8">
5  <style> /* set the CSS */
7
9   body { font: 12px Arial;}
11
13  path {
15    stroke: steelblue;
16    stroke-width: 2;
17    fill: none;
18 }
19
21
23 .axis path,
24 .axis line {
25   fill: none;
26   stroke: grey;
27   stroke-width: 1;
28   shape-rendering: crispEdges;
29 }
30
32
34 rect.bordered {
35   stroke: #E6E6E6;
36   stroke-width: 2px;
37 }
38
39
41 text.mono {
42   font-size: 6pt;
43   font-family: Consolas, courier;
44   fill: #aaa;
45 }
46
48 text.axis-workweek {
49   fill: #000;
50 }
51
53 text.axis-worktime {
```

32 A. APPENDIX A

```
39         fill: #000;
      }
      .RdYlGn .q0-11{ fill:rgb(165,0,38)}
      .RdYlGn .q1-11{ fill:rgb(215,48,39)}
      .RdYlGn .q2-11{ fill:rgb(244,109,67)}
      .RdYlGn .q3-11{ fill:rgb(253,174,97)}
      .RdYlGn .q4-11{ fill:rgb(254,224,139)}
      .RdYlGn .q5-11{ fill:rgb(255,255,191)}
      .RdYlGn .q6-11{ fill:rgb(217,239,139)}
      .RdYlGn .q7-11{ fill:rgb(166,217,106)}
      .RdYlGn .q8-11{ fill:rgb(102,189,99)}
      .RdYlGn .q9-11{ fill:rgb(26,152,80)}
      .RdYlGn .q10-11{ fill:rgb(0,104,55)}

53 .header {
  height:50px;
  background:#F0F0F0;
  border:1px solid #CCC;
  width:960px;
  margin:0px auto;
}
</style>
<body>

63 <!— <div id="option">
  <input name="updateButton"
    type="button"
    value="Previous"
    onclick="previousData() " />
</div>
69
<div id="option">
  <input name="updateButton"
    type="button"
    value="Next"
    onclick="nextData() " />
</div> —>
75

77 <!— load the d3.js library —>
<script src="http://d3js.org/d3.v3.min.js"<!--/script--&gt;
&lt;script type='text/javascript' src='knockout-min.js'><!--/script--&gt;

81 &lt;div style="font-size: 50px; text-align: center; margin-top: 20px; margin
  -bottom: 20px" data-bind="text: currentDate() "><!--/div--&gt;
&lt;div id="year" style="margin-top=20px"--&gt;/div&gt;
83 &lt;div id="area1"--&gt;/div&gt;
&lt;div id="area2" style="margin-bottom:20px"--&gt;/div&gt;
85
&lt;script&gt;
87   function AppViewModel(){
89     this.currentDay = ko.observable(1);</pre>
```

```

availableCountries = ko.observableArray(['129.432.123.3',
91   '129.987.123.422',
   '987.654.432.123','456.423.0.124','126.674.234.543','129.432.123.3',
   '129.987.123.422',
   '987.654.432.123','456.423.0.124','126.674.234.543','129.432.123.3',
   '129.987.123.422',
   '987.654.432.123','456.423.0.124','126.674.234.543']);
availablePorts = ko.observableArray
([['80','8080','8000','7000','9000','8888','80','8080','8000','7000','9000','8888']);

this.chosenIp = ko.observable(availableCountries()[0]);
this.chosenPort = ko.observable(availablePorts()[0]);
this.currentYear = ko.observable(2008);
this.currentMonth = ko.observable(1);
this.currentDate = ko.observable('2008-01-01');
97 }

99 var parseDate = d3.time.format("%d-%b-%Y").parse;
101 ko.applyBindings(AppViewModel);

103 // ** Update data section (Called from the onclick)

105
106 function firstGraph(){
107   var width = 2000,
108     height = 250,
109     cellSize = 35;
110     date = "01-01-2012" // cell size
111
112   var percent = d3.format(".1%"),
113     format = d3.time.format("%Y-%m-%d");
114
115   var color = d3.scale.quantize()
116     .domain([- .5, .5])
117     .range(d3.range(11).map(function(d) { return "q" + d + "-11"; }));
118
119   var svg = d3.select("#year").selectAll("svg")
120     .data(d3.range(2008, 2009))
121     .enter().append("svg")
122     .attr("width", width)
123     .attr("height", height)
124     .attr("class", "RdYlGn")
125     .append("g")
126     .attr("transform", "translate(" + ((width - cellSize * 53) / 2) +
127       " , " + (height - cellSize * 7 - 1) + ")");
128
129   svg.append("text")
130     .attr("transform", "translate(-6," + cellSize * 3.5 + ") rotate(-90")
131     .style("text-anchor", "middle")
132     .text(function(d) { return d; });

```

```

133 var rect = svg.selectAll(".day")
134   .data(function(d) { return d3.time.days(new Date(d, 0, 1), new Date
135     (d + 1, 0, 1)); })
136   .enter().append("rect")
137   .attr("class", "day")
138   .attr("width", cellSize)
139   .attr("height", cellSize)
140   .attr("x", function(d) { return d3.time.weekOfYear(d) * cellSize; })
141   .attr("y", function(d) { return d.getDay() * cellSize; })
142   .on("click", function(d){
143     heatmapChart("heatmap/" + d + ".csv");
144     currentDate(d);
145     console.log(currentDate());
146     updateGraph();
147   })
148   .datum(format);

149 rect.append("title")
150   .text(function(d) { return d; });

151 svg.selectAll(".month")
152   .data(function(d) { return d3.time.months(new Date(d, 0, 1), new
153     Date(d + 1, 0, 1)); })
154   .enter().append("path")
155   .attr("class", "month")
156   .attr("d", monthPath);

157 d3.csv("year/dji.csv", function(error, csv) {
158   if (error) throw error;

159   var data = d3.nest()
160     .key(function(d) { return d.Date; })
161     .rollup(function(d) { return d[0].High; })
162     .map(csv);

163   rect.filter(function(d) { return d in data; })
164     .attr("class", function(d) { return "day" + color(data[d] % 11); })
165     .select("title")
166       .text(function(d) { return d + ": " + (data[d]); });
167 });

168 function monthPath(t0) {
169   var t1 = new Date(t0.getFullYear(), t0.getMonth() + 1, 0),
170     d0 = t0.getDay(), w0 = d3.time.weekOfYear(t0),
171     d1 = t1.getDay(), w1 = d3.time.weekOfYear(t1);
172   return "M" + (w0 + 1) * cellSize + "," + d0 * cellSize
173     + "H" + w0 * cellSize + "V" + 7 * cellSize
174     + "H" + w1 * cellSize + "V" + (d1 + 1) * cellSize
175     + "H" + (w1 + 1) * cellSize + "V" + 0

```

```

    + "H" + (w0 + 1) * cellSize + "Z";
181 }

183 d3.select(self.frameElement).style("height", "2910px");

185 var margin = { top: 50, right: 0, bottom: 10, left: 60 },
187     width = 2000 - margin.left - margin.right,
188     height = 650 - margin.top - margin.bottom,
189     gridSize = Math.floor(width / 24),
190     legendElementWidth = gridSize,
191     buckets = 9,
192     colors = ["#ffffd9", "#edf8b1", "#c7e9b4", "#7fcdbb", "#41b6c4", "#1d91c0", "#225ea8", "#253494", "#081d58"], // alternatively
193     colorbrewer.YlGnBu[9]
194     days = availablePorts(),
195     times = availableCountries(),
196     datasets = ["heatmap/data.tsv", "heatmap/data2.tsv"];

197 var svg = d3.select("#area1")
198     .append("svg")
199     .attr("width", width + margin.left + margin.right)
200     .attr("height", height + margin.top + margin.bottom)
201     .append("g")
202     .attr("transform", "translate(" + margin.left + ", " + margin.
203     top + ")");
204
205 var dayLabels = svg.selectAll(".dayLabel")
206     .data(days)
207     .enter().append("text")
208     .text(function(d) { return d; })
209     .attr("x", 0)
210     .attr("y", function(d, i) { return i * gridSize; })
211     .style("text-anchor", "end")
212     .attr("transform", "translate(-6," + gridSize / 1.5 + ")")
213     .attr("class", function(d, i) { return ((i >= 0 && i <= 4)
? "dayLabel mono axis axis-workweek" : "dayLabel mono axis"); });
214
215 var timeLabels = svg.selectAll(".timeLabel")
216     .data(times)
217     .enter().append("text")
218     .text(function(d) { return d; })
219     .attr("x", function(d, i) { return i * gridSize; })
220     .attr("y", 0)
221     .style("text-anchor", "middle")
222     .attr("transform", "translate(" + gridSize / 2 + ", -6)")
223     .attr("class", function(d, i) { return ((i >= 7 && i <= 16)
? "timeLabel mono axis axis-worktime" : "timeLabel mono axis"); });
224
225 var heatmapChart = function(csvFile) {
226     d3.tsv(csvFile,

```

```

227     function(d) {
228         return {
229             day: +d.day,
230             hour: +d.hour,
231             value: +d.value,
232             ip: d.ip,
233             port: d.port
234         };
235     },
236     function(error, data) {
237         var colorScale = d3.scale.quantile()
238             .domain([0, buckets - 1, d3.max(data, function(d) {
239                 return d.value; }))])
240             .range(colors);
241
242         var cards = svg.selectAll(".hour")
243             .data(data, function(d) { return d.day+':'+d.hour; });
244
245         cards.append("title");
246
247         cards.enter().append("rect")
248             .attr("x", function(d) { return (d.hour - 1) * gridSize;
249 })
250             .attr("y", function(d) { return (d.day - 1) * gridSize;
251 })
252             .attr("rx", 4)
253             .attr("ry", 4)
254             .attr("class", "hour bordered")
255             .attr("width", gridSize)
256             .attr("height", gridSize)
257             .style("fill", colors[0])
258             .on("click", function(d){
259                 chosenIp(d.ip);
260                 chosenPort(d.port)
261                 console.log(d.value+" "+chosenIp()+" "+chosenPort())
262             ;
263                 updateGraph();
264             });
265
266         cards.transition().duration(1000)
267             .style("fill", function(d) { return colorScale(d.value);
268 });
269
270         cards.select("title").text(function(d) { return d.ip +'and' +
271             d.port+':'+d.value; });
272
273         cards.exit().remove();
274
275     });
276
277     heatmapChart('/heatmap/2008-01-01.csv');

```

```

273     var datasetpicker = d3.select("#dataset-picker").selectAll(".dataset-button")
274         .data(datasets);
275
276     datasetpicker.enter()
277         .append("input")
278         .attr("value", function(d){ return "Dataset " + d })
279         .attr("type", "button")
280         .attr("class", "dataset-button")
281         .on("click", function(d) {
282             updateGraph();
283         });
284
285 function updateGraph(){
286     d3.select("#area2").selectAll("svg").remove();
287
288     var margin = {top: 30, right: 20, bottom: 30, left: 70},
289     width = 2000 - margin.left - margin.right,
290     height = 500 - margin.top - margin.bottom;
291
292     // Parse the date / time
293
294     // Set the ranges
295     var x = d3.time.scale().range([0, width]);
296     var y = d3.scale.linear().range([height, 0]);
297
298     // Define the axes
299     var xAxis = d3.svg.axis().scale(x)
300         .orient("bottom").ticks(5);
301
302     var yAxis = d3.svg.axis().scale(y)
303         .orient("left").ticks(5);
304
305     // Define the line
306     var valueline = d3.svg.line()
307         .x(function(d) { return x(d.date); })
308         .y(function(d) { return y(d.close); });
309
310     // Adds the svg canvas
311     var svg = d3.select("#area2")
312         .append("svg")
313             .attr("width", width + margin.left + margin.right)
314             .attr("height", height + margin.top + margin.bottom)
315         .append("g")
316             .attr("transform",
317                 "translate(" + margin.left + ", " + margin.top + ")");
318
319     // Get the data
320     d3.csv("chart/" + currentDate() + '_' + chosenIp() + '_' + chosenPort() + '.csv',
321             function(error, data) {

```

```

323     data.forEach(function(d) {
324         d.date = +d.date;
325         d.close = +d.close;
326     });
327
328     // Scale the range of the data
329     x.domain(d3.extent(data, function(d) { return d.date; }));
330     y.domain([0, d3.max(data, function(d) { return d.close; })]);
331
332     // Add the valueline path.
333     svg.append("path")
334         .attr("class", "line")
335         .attr("d", valueline(data));
336
337     // Add the X Axis
338     svg.append("g")
339         .attr("class", "x axis")
340         .attr("transform", "translate(0," + height + ")");
341         .call(xAxis);
342
343     // Add the Y Axis
344     svg.append("g")
345         .attr("class", "y axis")
346         .call(yAxis);
347
348     });
349
350     updateGraph()
351 };
352
353 function changeIpOrPort() {
354
355     // Get the data again
356     d3.csv("data1.csv", function(error, data) {
357         data.forEach(function(d) {
358             d.date = parseDate(d.date);
359             d.close = +d.close;
360         });
361
362         // Scale the range of the data again
363         x.domain(d3.extent(data, function(d) { return d.date; }));
364         y.domain([0, d3.max(data, function(d) { return d.close; })]);
365
366         // Select the section we want to apply our changes to
367         var svg = d3.select("#area2").transition();
368
369         // Make the changes
370         svg.select(".line") // change the line
371             .duration(750)
372             .attr("d", valueline(data));
373         svg.select(".x.axis") // change the x axis

```

```
375     .duration(750)
376     .call(xAxis);
377     svg.select(".y.axis") // change the y axis
378     .duration(750)
379     .call(yAxis);
380   });
381 }
382 firstGraph();
383
384 </script>
385 </body>
```


Appendix B

Appendix B

Appendix B contains the scripts used to create .csv files from all the data made available from UNINETT.

A script that creates .csv files for every nfcapd file in a day. This script is run by another short script that runs it 31 times for each day.

```
102 #!/bin/bash
103 mkdir /home/eeglarose/flowtest/2012_02/$(printf "%02d" $1)
104 clock_converter(){
105     if [ $(($1 % 100)) -gt 59 ]; then
106         return $(($1 % 100))
107     fi
108     if [ $(($1 % 100)) -lt 60 ]; then
109         echo $(printf "%04d" $1)
110         return $(($1 % 100))
111     fi
112 }
113
114 for (( c=0; c<=2355; c += 5 ))
115 do
116     ndump -r $(printf "%02d" $1)/nfcapd.201202$(printf "%02d" $1)$(
117         clock_converter $c ) -n 10 -s srcip -o csv > /home/eeglarose/
118         flowtest/2012_02/$(printf "%02d" $1)/$(clock_converter $c ).csv
119 done
```

Listing B.1: Creates .csv files for every nfcapd file in a day

A script that fetches the total amount of flows for each day and creates a file with the values.

```
118
119
120 clock_converter(){
121     if [ $(($1 % 100)) -gt 59 ]; then
122         return $(($1 % 100))
123     fi
```

```

124 if [ $($1 % 100) -lt 60 ]; then
125   echo $(printf "%04d" $1)
126   return $($1 % 100)
127 fi
128 }

130 for (( c=0; c<=2355; c += 5 ))
131 do
132   total_file=$(awk -F',' 'NR == 15 { print $1}' /home/eeglarse/
133   flowtest/2012_02/$(printf "%02d" $1)/$( clock_converter $c ).csv)
134   echo $total_file >> testfile2$1.csv
135 done
136
137 awk '{s+=$1} END {print s}' testfile2$1.csv >>datefile2.csv

```

Listing B.2: Total amount of flows for each day

A script that finds the top 10 used IP-adresses for each day.

```

140 nfdump -R /data/netflow/oslo_gw/2012/01/$(printf "%02d" $1)/nfcapd
      .201201$(printf "%02d" $1)0000:nfcapd.201201$(printf "%02d" $1)2355
      -n 10 -s dstip -o csv > /home/eeglarse/flowtest/top10/$(printf "%02d" $1).csv

```

Listing B.3: Top 10 used IP-adresses for each day

A script that creates a list the top 10 most popular ports, based on the 10 most popular IP-addresses.

```

1 ip_string=''
2 ip=$(awk -F',' 'NR == 2 { print $5}' /home/eeglarse/flowtest/top10/$(
      printf "%02d" $1).csv)
ip_string+='dst ip '$ip' or '
4 ip=$(awk -F',' 'NR == 3 { print $5}' /home/eeglarse/flowtest/top10/$(
      printf "%02d" $1).csv)
ip_string+='dst ip '$ip' or '
6 ip=$(awk -F',' 'NR == 4 { print $5}' /home/eeglarse/flowtest/top10/$(
      printf "%02d" $1).csv)
ip_string+='dst ip '$ip' or '
8 ip=$(awk -F',' 'NR == 5 { print $5}' /home/eeglarse/flowtest/top10/$(
      printf "%02d" $1).csv)
ip_string+='dst ip '$ip' or '
10 ip=$(awk -F',' 'NR == 6 { print $5}' /home/eeglarse/flowtest/top10/$(
      printf "%02d" $1).csv)
ip_string+='dst ip '$ip' or '
12 ip=$(awk -F',' 'NR == 7 { print $5}' /home/eeglarse/flowtest/top10/$(
      printf "%02d" $1).csv)
ip_string+='dst ip '$ip' or '
14 ip=$(awk -F',' 'NR == 8 { print $5}' /home/eeglarse/flowtest/top10/$(
      printf "%02d" $1).csv)

```

```

1 ip_string+='dst_ip '$ip' or '
16 ip=$(awk -F',' 'NR == 9 { print $5}' /home/eeglarse/flowtest/top10/$(
    printf "%02d" $1).csv)
    ip_string+='dst_ip '$ip' or '
18 ip=$(awk -F',' 'NR == 10 { print $5}' /home/eeglarse/flowtest/top10/$(
    printf "%02d" $1).csv)
    ip_string+='dst_ip '$ip' or '
20 ip=$(awk -F',' 'NR == 11 { print $5}' /home/eeglarse/flowtest/top10/$(
    printf "%02d" $1).csv)
    ip_string+='dst_ip '$ip
22

24 nfdump -R /data/netflow/oslo_gw/2012/01/$(printf "%02d" $1)/nfcapd
    .201201$(printf "%02d" $1)0000:nfcapd.201201$(printf "%02d" $1)2355
        -n 10 -s dstport $iplist -o csv > /home/eeglarse/flowtest/top10/
            top10port/$(printf "%02d" $1).csv

```

A script that uses the 10 most popular IP-adresses and their corresponding ports to find the number of flows sent to each port on each IP-address.

```

2 #!/bin/bash
4 for (( i = 1; i < 31; i++ )); do
5     iplist=()
6     ip=$(awk -F',' 'NR == 2 { print $5}' /home/eeglarse/flowtest/top10/$(
7         printf "%02d" $i).csv)
8     iplist[0]=$ip
9     ip2=$(awk -F',' 'NR == 3 { print $5}' /home/eeglarse/flowtest/top10/$(
10        printf "%02d" $i).csv)
11    iplist[1]=$ip2
12    ip=$(awk -F',' 'NR == 4 { print $5}' /home/eeglarse/flowtest/top10/$(
13        printf "%02d" $i).csv)
14    iplist[2]=$ip
15    ip=$(awk -F',' 'NR == 5 { print $5}' /home/eeglarse/flowtest/top10/$(
16        printf "%02d" $i).csv)
17    iplist[3]=$ip
18    ip=$(awk -F',' 'NR == 6 { print $5}' /home/eeglarse/flowtest/top10/$(
19        printf "%02d" $i).csv)
20    iplist[4]=$ip
21    ip=$(awk -F',' 'NR == 7 { print $5}' /home/eeglarse/flowtest/top10/$(
22        printf "%02d" $i).csv)
23    iplist[5]=$ip
24    ip=$(awk -F',' 'NR == 8 { print $5}' /home/eeglarse/flowtest/top10/$(
25        printf "%02d" $i).csv)
26    iplist[6]=$ip
27    ip=$(awk -F',' 'NR == 9 { print $5}' /home/eeglarse/flowtest/top10/$(
28        printf "%02d" $i).csv)
29    iplist[7]=$ip
30    ip=$(awk -F',' 'NR == 10 { print $5}' /home/eeglarse/flowtest/top10/$(
31        printf "%02d" $i).csv)
32    iplist[8]=$ip

```

```

ip=$(awk -F' , ' 'NR == 11 { print $5}' /home/eeglarse/flowtest/top10/$
    (printf "%02d" $i).csv)
iplist[9]=$ip
portlist=()
ip=$(awk -F' , ' 'NR == 2 { print $5}' /home/eeglarse/flowtest/top10/
    top10port/$(printf "%02d" $i).csv)
portlist[0]=$ip
ip=$(awk -F' , ' 'NR == 3 { print $5}' /home/eeglarse/flowtest/top10/
    top10port/$(printf "%02d" $i).csv)
portlist[1]=$ip
ip=$(awk -F' , ' 'NR == 4 { print $5}' /home/eeglarse/flowtest/top10/
    top10port/$(printf "%02d" $i).csv)
portlist[2]=$ip
ip=$(awk -F' , ' 'NR == 5 { print $5}' /home/eeglarse/flowtest/top10/
    top10port/$(printf "%02d" $i).csv)
portlist[3]=$ip
ip=$(awk -F' , ' 'NR == 6 { print $5}' /home/eeglarse/flowtest/top10/
    top10port/$(printf "%02d" $i).csv)
portlist[4]=$ip
ip=$(awk -F' , ' 'NR == 7 { print $5}' /home/eeglarse/flowtest/top10/
    top10port/$(printf "%02d" $i).csv)
portlist[5]=$ip
ip=$(awk -F' , ' 'NR == 8 { print $5}' /home/eeglarse/flowtest/top10/
    top10port/$(printf "%02d" $i).csv)
portlist[6]=$ip
ip=$(awk -F' , ' 'NR == 9 { print $5}' /home/eeglarse/flowtest/top10/
    top10port/$(printf "%02d" $i).csv)
portlist[7]=$ip
ip=$(awk -F' , ' 'NR == 10 { print $5}' /home/eeglarse/flowtest/top10/
    top10port/$(printf "%02d" $i).csv)
portlist[8]=$ip
ip=$(awk -F' , ' 'NR == 11 { print $5}' /home/eeglarse/flowtest/top10/
    top10port/$(printf "%02d" $i).csv)
portlist[9]=$ip
for (( s = 0; s < 10; s++ )); do
    for (( j = 0; j < 10; j++ )); do
        $(nfdump -R /data/netflow/oslo_gw/2012/01/$(printf "%02d" $i)/
            nfcapd.201201$(printf "%02d" $i)0000:nfcapd.201201$(printf "%02d"
            $i)2355 -n 10 -s dstport -o csv 'dst ip ${iplist[$s]} and dst port
            ${portlist[$j]}' -o csv
        done
    done
done
done

```